



# AI USAGE POLICY

---

Version 1.1

Effective Date: June 7, 2026

Last Reviewed: June 7, 2026

Policy Owner: Rebecca Younes, President

## Table of Contents

|                                                 |   |
|-------------------------------------------------|---|
| 1. Purpose and Scope .....                      | 3 |
| 2. Guiding Principles .....                     | 3 |
| 3. How We Use AI to Support Our Work.....       | 3 |
| 4. Company-Managed Tools and Default Usage..... | 4 |
| 4.1 Why Use Company-Managed Tools.....          | 4 |
| 4.2 Company-Managed Tools.....                  | 4 |
| 4.3 Use and Visibility of Company Tools .....   | 5 |
| 5. Recommended and Discouraged Uses of AI ..... | 5 |
| 5.1 Recommended Uses.....                       | 5 |
| 5.2 Use with Caution.....                       | 6 |
| 5.2.1 AI-Assisted Technical Tasks .....         | 6 |
| 5.3 Discouraged Uses.....                       | 7 |
| 5.3.1 Ontology Design .....                     | 8 |
| 5.3.2 Independent Thinking .....                | 8 |
| 6. Use of Non-Company-Managed Tools .....       | 8 |
| 6.1 Permitted Uses .....                        | 8 |
| 6.2 Prohibited Uses .....                       | 8 |

|                                                                 |    |
|-----------------------------------------------------------------|----|
| 6.3 Personal Accounts.....                                      | 9  |
| 7. Data Protection and Confidentiality.....                     | 9  |
| 7.1 General Data Handling Principles .....                      | 9  |
| 7.2 Sensitive Information.....                                  | 9  |
| 7.3 Highly Sensitive Information.....                           | 10 |
| 7.4 Use of Non-Company-Managed Tools .....                      | 10 |
| 7.5 Guiding Rules .....                                         | 10 |
| 8. Responsibility and Ownership .....                           | 11 |
| 9. AI Tools with System or Resource Access.....                 | 11 |
| 9.1 General Rules.....                                          | 12 |
| 9.2 Isolation for AI Tools with System or Resource Access ..... | 12 |
| 9.3 Technical Guidelines for AI-Assisted Coding.....            | 13 |
| 9.3.1 General .....                                             | 13 |
| 9.3.2 Prompt Construction .....                                 | 13 |
| 9.3.3 GitHub Interactions.....                                  | 13 |
| 9.4 Other System Integrations.....                              | 14 |
| 10. Client Disclosure and Approval.....                         | 14 |
| 11. When in Doubt.....                                          | 14 |
| 12. Acknowledgement of Policy .....                             | 15 |

# 1. Purpose and Scope

This policy establishes guidelines for the responsible, secure, and effective use of artificial intelligence (AI) tools within Semantic Arts. It is intended to enable productive use of AI while protecting client confidentiality, safeguarding company intellectual property, and maintaining high standards of professional work.

This policy applies to all employees and interns, and any contractors or third parties using AI tools in connection with company or client work (collectively referred to as “employees”).

For the purposes of this policy, “AI tools” include large language models, code assistants, and AI-powered applications used for generating, analyzing, or modifying content.

---

## 2. Guiding Principles

Semantic Arts employees are expected to use AI tools in accordance with the following principles:

- Protect client confidentiality at all times.
  - Protect company intellectual property and methodologies.
  - Use AI to support—not replace—professional judgment and domain expertise.
  - Maintain full responsibility for all work products.
  - Be transparent about AI use, especially in client work.
- 

## 3. How We Use AI to Support Our Work

We use AI to enhance the quality, depth, and effectiveness of our work, not simply to complete the same work more quickly. Productivity gains from AI are expected to translate into greater output, broader impact, and improved outcomes for our clients.

The demand for high-quality semantic modeling, data integration, and knowledge graph development far exceeds our current capacity. AI enables us to address more of that demand by accelerating implementation, analysis, and other supporting activities, allowing us to deliver more value within the same time and resource constraints and devote more attention to semantic modeling, where human judgment, insight, and domain expertise provide the greatest value.

Employees are expected to use AI to support their professional judgment and increase the value they deliver, not to reduce their level of contribution. This approach reinforces our company values of “Create value” and “Share equitably”: we use AI to increase the value we provide our clients rather than to shift the balance of revenue away from the consultants who earn it.

We encourage employees to use AI not only to complete individual tasks, but also to create reusable assets that increase the effectiveness of future work. Examples include prompt templates, documented workflows, shared knowledge resources, automation scripts, and other tools that can be reused across projects and teams.

---

## 4. Company-Managed Tools and Default Usage

Using company-managed AI tools is the default for work-related AI usage.

---

### 4.1 Why Use Company-Managed Tools

Using company-managed AI tools increases the likelihood of safe and compliant usage, as these tools are configured with organization-level security settings. Organization-level settings take precedence over individual configurations.

Employees are encouraged to report any limitations or obstacles encountered so they can be evaluated and, where appropriate, addressed.

---

### 4.2 Company-Managed Tools

Semantic Arts currently provides the following AI tools:

- ChatGPT, including Codex (company account)
- Claude, including Claude Code and CoWork (company account)

Only these tools may be used for any work involving:

- Client data
- Internal intellectual property
- Deliverables prepared for clients

Employees may request access to one or both of the company-managed AI tools.

Semantic Arts maintains flexibility in its tool selection and evaluates available options on an ongoing basis to ensure that employees have access to effective and appropriate AI tools. Employee feedback based on practical experience is encouraged and welcomed.

The company will also provide opportunities for knowledge sharing, presentations, and structured learning. Participation in these activities is strongly encouraged and, in some cases, may be required.

The company will maintain a repository of shared resources, including prompt templates, use case documentation, and other materials that support a growing firm-wide AI knowledge base. Employees are expected to contribute to and maintain these resources as part of their work.

---

## 4.3 Use and Visibility of Company Tools

Use of company-provided AI tools may be subject to administrative visibility, including usage metrics and, in some cases, interaction data. Employees should not assume that prompts, outputs, or uploaded content are private when using these tools.

Company-provided AI tools are intended for work-related use and should be treated similarly to other company systems, such as email or document storage. Personal use should be limited and must not involve sensitive, confidential, or inappropriate content.

---

# 5. Recommended and Discouraged Uses of AI

AI tools can be valuable assistants when used appropriately. The following guidance outlines where AI is most effective and where its use should be limited.

---

## 5.1 Recommended Uses

AI tools are well-suited for:

- Drafting and refining definitions, descriptions, and documentation
- Correcting grammar and improving clarity and writing style
- Brainstorming terminology and naming
- Summarizing non-sensitive content
- Generating examples to illustrate and communicate concepts
- Presenting options and alternatives
- Supporting research on publicly available information

- Producing graphics and visual content for presentations and marketing materials
- Generating outputs that can be studied to build understanding and expertise

These activities are most effective when performed interactively and through cycles of refinement rather than as one-time requests.

---

## 5.2 Use with Caution

AI may be used with caution for:

- Reviewing or suggesting improvements to existing work
- Identifying potential gaps or inconsistencies
- Proposing alternative approaches for consideration

Outputs in these cases must be treated as suggestions and carefully evaluated.

Certain use cases carry additional risk. In particular, extracting or processing information from financial documents (e.g., receipts, invoices, or billing emails) should be approached with care. These materials often contain sensitive information, and errors in extraction can lead to incorrect financial or operational outcomes.

---

### 5.2.1 AI-Assisted Technical Tasks

AI tools may be used to assist with technical tasks such as:

- Coding and scripting, including code documentation
- Assisting with code review and analysis
- Data mapping and transformation
- Writing queries
- Generating sample data
- Developing spreadsheets, formulas, and models
- Graph discovery

Data mapping tasks vary in risk depending on how AI is used. Generating mappings based on defined patterns is generally lower risk, while generating or refining mapping patterns is higher risk and more akin to semantic modeling.

Use of AI for any of the above tasks should be approached with caution:

- Outputs must be reviewed, tested, and validated before use.
- Results must align with company standards, architecture, design patterns, and intended semantics.

- AI-generated outputs may contain errors, inefficiencies, or incorrect assumptions, including overlooked edge cases or context-specific meaning.
- Use of AI should support implementation and productivity, not replace understanding of the underlying logic, data, or models.

Employees remain fully responsible for the quality, correctness, and maintainability of all technical outputs.

In addition to validation, employees are encouraged to review and study AI-generated outputs to understand how solutions are constructed and build professional skills.

AI tools may also be used to generate or assemble applications, workflows, or internal tools, including by employees without a technical background. This can be an effective use of AI but requires additional care.

When using AI in this way:

- Ensure that the functionality of the application is understood at a practical level (what it does, what data it uses, and how it behaves).
- Validate that outputs are appropriate for the intended use.
- Do not use such tools for sensitive, client-facing, or business-critical processes without appropriate review.
- Seek assistance when needed to evaluate risks, limitations, or technical concerns.
- Applications that store or process sensitive data must not be created or deployed without appropriate review.

Use of AI to generate applications does not remove responsibility for the behavior, accuracy, or impact of those applications.

---

## 5.3 Discouraged Uses

AI tools should not be relied upon for:

- Designing ontologies, schemas, or data models
- Defining core domain structures or relationships
- Making decisions requiring deep domain expertise or client-specific context
- Making business, strategic, or organizational decisions

AI may be used to support business and strategic decision-making through research, analysis, and generation of alternatives, but responsibility for decisions remains with those responsible for making them.

### 5.3.1 Ontology Design

We have tested the use of AI tools for generating ontologies and found that, even when provided with our design patterns and best practices, the outputs lack the conceptual depth and insight required for high-quality modeling.

Effective ontology design requires human judgment to align models with business concepts and use cases in a way that is simple, elegant, and extensible.

---

### 5.3.2 Independent Thinking

AI tools should not be used as a substitute for initial problem framing and independent thinking.

Employees are expected to develop a basic understanding of the problem, context, and objectives before using AI tools. Starting with AI-generated answers without this foundation can lead to shallow, incorrect, or misaligned results.

AI is most effective when used to refine, extend, or challenge an initial line of thought, not to replace it.

---

## 6. Use of Non-Company-Managed Tools

---

### 6.1 Permitted Uses

Employees may use non-company-managed AI tools without prior approval only for the following purposes:

- Learning and skill development
  - Brainstorming and ideation
  - General productivity tasks
  - Work involving non-sensitive, publicly available information
- 

### 6.2 Prohibited Uses

Non-company-managed tools must not be used for:

- Client data of any kind



- Client deliverables and work products
  - Internal company documents or intellectual property
  - Proprietary code or technical assets
  - Any information covered by NDA or confidentiality obligations
- 

## 6.3 Personal Accounts

Personal AI accounts may not be used for:

- Client-related work
  - Internal company work
  - Any task involving confidential or proprietary information
- 

# 7. Data Protection and Confidentiality

The following data protection principles apply to the use of all AI tools, whether company-managed or non-company-managed.

---

## 7.1 General Data Handling Principles

Employees must exercise care when entering information into any AI tool.

- Only enter sensitive or confidential information when necessary and when using company-managed tools.
  - Consider how information may be stored, processed, or reused.
  - Ensure compliance with client obligations and company policies.
- 

## 7.2 Sensitive Information

The following are considered sensitive and must be protected:

- Client data and business information
- Ontologies, schemas, taxonomies, and semantic models
- Internal methodologies, frameworks, and processes
- Proprietary code and technical assets, including the contents of Semantic Arts' private GitHub repositories
- Any information covered by NDA or confidentiality obligations

---

## 7.3 Highly Sensitive Information

Certain types of information require heightened caution, even when using approved AI tools:

- Personally identifiable information (PII)
- Financial data (company or client)
- Any regulated or highly confidential information, such as HIPAA-regulated data

Additional caution should be used when working with large amounts of data, as they may increase the risk of unintended exposure or misuse.

When using approved AI tools:

- Limit information to what is strictly necessary and avoid entering bulk data or full records.
- Prefer anonymized, synthesized, summarized, or generalized information where possible.

With client approval, anonymized or synthesized data may be entered into an approved AI tool. Anonymization must be thorough — replacing names with generic placeholders (e.g., "Client A", "User 1") and removing all unique identifiers.

Highly sensitive information must never be entered into non-company-managed AI tools.

---

## 7.4 Use of Non-Company-Managed Tools

When using non-company-managed AI tools, employees must assume that any information entered:

- May be stored
- May be reviewed by the provider
- May be used for model training

Accordingly, sensitive or highly sensitive information must not be entered into non-company-managed tools under any circumstances.

---

## 7.5 Guiding Rules

If you would not share the information in a public forum, do not enter it into an AI tool unless using an approved tool in accordance with this policy.

Use the minimum amount of sensitive information necessary to accomplish the task.

---

## 8. Responsibility and Ownership

Employees are fully responsible for all work products they produce, regardless of whether AI tools were used.

- AI-generated content must be reviewed, validated, and refined before inclusion in deliverables.
- Outputs must meet company and client standards for accuracy and quality.
- Errors or omissions resulting from AI use remain the responsibility of the employee.

AI tools are assistants and do not replace professional accountability.

Work products created with the assistance of AI tools are considered company work product and are subject to the same ownership and intellectual property policies as all other work. Use of AI tools does not alter ownership or responsibility for deliverables.

---

## 9. AI Tools with System or Resource Access

This section applies to AI tools that can access or act upon files, repositories, applications, services, credentials, or other resources beyond the AI provider's own environment. This covers both:

- **Local AI tools** running on the user's device (installed applications, CLI utilities, IDE integrations), which act through the machine's file system, network access, and installed software.
- **Cloud or remote AI tools** that run elsewhere but act through the user's credentials or authorized account connections—OAuth tokens, API keys, SSO sessions, or stored passwords—to perform actions on services such as email, social platforms, code hosts, or other SaaS applications.

Examples of such tools include Claude Code and OpenAI Codex, which can modify files, execute code, and interact with code repositories; Claude Cowork, which automates desktop file and task operations; and workflow-integration tools that act on connected SaaS accounts.

These tools require safeguards beyond those needed for prompt-only AI use. These safeguards are described below, and they apply regardless of where the tool runs. A local tool is not inherently contained—a user's device typically holds SSH keys, cached SSO tokens, and active VPN or browser sessions that bridge into company systems. A cloud or remote tool is not inherently lower-risk—it acts with the user's credentials on connected services, and the resulting actions

(emails sent, posts published, code committed, data deleted) are attributable to the user and often difficult to reverse.

---

## 9.1 General Rules

- Use company-controlled AI accounts only.
  - Limit access to the minimum necessary scope.
  - Maintain human oversight of outputs and actions, not just final results.
  - Do not allow AI tools to access or operate directly on sensitive, production, or business-critical systems.
  - Configure AI integrations (including connectors and MCP servers) with the minimum permissions necessary.
- 

## 9.2 Isolation for AI Tools with System or Resource Access

AI tools that execute code or access systems should be run in isolated environments (e.g., containers, virtual machines, or even separate hardware) unless there is a clear and justified reason not to do so. A VM is safer than containers because it creates a stronger isolation boundary. Separate hardware offers the strongest isolation but is rarely necessary.

While isolation reduces risk, it does not eliminate it. Systems should be designed and used with the assumption that boundaries may be imperfect. In particular, isolation environments should:

- Use least-privilege access.
- Restrict network, file system, and system access within the execution environment to resources required for the intended operations.
- Provide sufficient logging to support audits and incident investigations.
- Do not store API keys, credentials, or sensitive configuration in locations accessible to AI tools.

Examples of environments that can provide sufficient isolation include:

- Containerized development environment (e.g., VS Code Dev Container, Docker-based sandbox) with local file access limited to the target repository and network access limited to GitHub, required package registries (PyPI, npm, etc.), and the AI vendor's API endpoints. (Container-based environments are appropriate for typical AI coding work; higher-risk tasks should use VM-level isolation.)
- Vendor-hosted VM that the user interacts with through a chat interface, with the tool's access to the user's systems limited to narrow, audited channels (e.g., a GitHub integration) and no direct access to the user's local files or credentials.

Note that some common measures intended to isolate environments do not provide the protections that people assume they do. For example, starting the tool in a specific directory does not actually prevent it from reading or writing in other directories outside it.

Further details about isolation environments will be published separately in a forthcoming handbook. The company will also provide tools to assist with installing, configuring, and using such solutions.

---

## 9.3 Technical Guidelines for AI-Assisted Coding

### 9.3.1 General

- Understand the problem and choose a solution before proceeding with coding. Use AI chat sessions to determine the tools, packages, and methods for a given problem.
- Ensure that features and bug fixes being implemented are covered by tests.
- Understand the changes being implemented. Do not commit a “black box” of code changes that you have not read and understood.

---

### 9.3.2 Prompt Construction

- Keep prompts unambiguous. Ambiguity invites unintentional changes and catastrophic results such as deletion of an entire database.
- Keep each prompt focused on a single problem. Do not specify too many tasks or problems in a single prompt.

---

### 9.3.3 GitHub Interactions

- AI systems should operate in a “propose, don’t execute” mode for actions that modify Semantic Arts GitHub repositories. Human approval is required before executing such actions.
  - The tool may open pull requests but must not have permission to push directly to protected branches or trigger production deployments. Changes proposed by AI tools should be reviewed and merged by a human before automated pipelines execute.
-

## 9.4 Other System Integrations

AI tools may be capable of connecting directly to company systems such as document storage (e.g., SharePoint), communication platforms (e.g., Slack, Teams), or other applications.

To protect client and company data:

- Non-company-managed AI tools must not be connected to company systems under any circumstances.
- Connections between AI tools and company systems may only be established using company-managed tools and configurations.
- Such integrations will be managed and approved at the company level, not by individual users.

Employees should not attempt to connect AI tools to company systems unless explicitly directed or enabled as part of an approved setup.

---

## 10. Client Disclosure and Approval

The use of AI tools in client-related work must be transparent.

- Project managers are responsible for ensuring that clients are informed of AI usage.
- Where required by contract or client preference, approval must be obtained before use.
- AI must not be used where clients prohibit its use or where disclosure has not occurred.

Employees should be prepared to explain how AI is used and how confidentiality and quality are maintained.

---

## 11. When in Doubt

If there is uncertainty about whether a tool or use case is appropriate:

- Default to company-managed tools.
  - Seek clarification from the policy owner.
-

## 12. Acknowledgement of Policy

Use of AI tools in connection with company or client work is considered part of normal work activity and subject to company policies.

Employees are expected to review this policy and any subsequent updates. Continued use of AI tools for work-related purposes constitutes acknowledgement of and agreement to comply with this policy.

Request access to one or more of the company-provided AI tools by sending an email to [office@semanticarts.com](mailto:office@semanticarts.com).

---

*This policy was developed with the assistance of AI tools in accordance with the guidelines outlined herein.*